


# Communication Functions

A **Gather** operation consists of every node sending a consistent amount of data to a single node.

## A few more MPI functions

`MPI_Gather` is a method for gathering information from multiple processes to one process.

# A few more MPI functions



```
MPI_Gather(sendbuf, sendcount, sendtype, recvbuf,  
           recvcount, root, comm)
```

- ▶ `void* sendbuf` – The start of the buffer to send.
- ▶ `int sendcount` – The number of items to send.
- ▶ `MPI_Datatype` – The type of data getting sent.
- ▶ `void* recvbuf` – The buffer to receive data. Only matters on the receiver.
- ▶ `int recvcount` – The number of items to receive **from each process**
- ▶ `int root` – The rank of the process that is collecting the data
- ▶ `MPI_Comm comm` – The communicator

```
if (rank==0){
    int mydata=-1;//root still sends data.
    //make room to collect
    int* recvbuf=malloc(total_procs*sizeof(int));
    MPI_Gather(&mydata,1,MPI_INT,recvbuf,1,MPI_INT,0,
        MPI_COMM_WORLD);
    for (i=0;i<total_procs;i++){
        printf("%d\n",recvbuf[i]);
    }
}
else{
    int mydata=rank*2;
    MPI_Gather(&mydata,1,MPI_INT,NULL,0,MPI_INT,0,
        MPI_COMM_WORLD);
}
```

```
[joe@eichholz-2 MA335]$ mpirun -np 4 a.out
```

```
-1
```

```
2
```


```
4
```

```
6
```

# Scatter

A **scatter operation** takes a large amount of data on a single node and distributes the data equally among the other nodes.

## A few more MPI functions





```
int MPI_Scatter(const void *sendbuf, int sendcount,  
    MPI_Datatype sendtype, void *recvbuf,  
    int recvcount, MPI_Datatype recvtype,  
    int root, MPI_Comm comm)
```

- ▶ `sendbuf` - The buffer to be send. Significant only to the sender.
- ▶ `sendcount` - The number of elements of `sendbuf` to send each processor
- ▶ `sendtype` - The datatype being sent.
- ▶ `recvbuf` - The receive buffer. Significant only on receivers.
- ▶ `recvcount` - The number of elements to receive.
- ▶ `recvtype` - The datatype to receive.
- ▶ `root` - The rank of the processor doing the sending.
- ▶ `comm` - The communicator over which this is happening.

MPI\_Allgather is just like MPI\_Gather, except that with Allgather the data shows up on every process instead of just one copy on the root.

```
if (rank==0){
    int mydata=-1;//root still sends data.
    int* recvbuf=new int[size]; //make room to collect
    MPI_Allgather(&mydata,1,MPI_INT,recvbuf,1,MPI_INT,
        MPI_COMM_WORLD);
}
else{
    int mydata=rank*2;
    int* recvbuf=new int[size];
    MPI_Allgather(&mydata,1,MPI_INT,recvbuf,1,MPI_INT,
        MPI_COMM_WORLD);
    for (int i=0;i<size;i++){
        cout<<"Rank: "<<rank<<" index: "<<i<<" "<<recvbuf[i]<
    }
}
```

```
[joe@eichholz-2 MA335]$ mpirun -np 3 a.out
```

```
Rank: 1index: 0 -1
```

```
Rank: 1index: 1 2
```

```
Rank: 1index: 2 4
```

```
Rank: 2index: 0 -1
```

```
Rank: 2index: 1 2
```

```
Rank: 2index: 2 4
```

```
MPI_Allgather(sendbuf, sendcount, sendtype, recvbuf, recvcount,  
recvtype, comm)
```

- ▶ Exactly the same as for MPI\_Gather, note that there is no root since everyone is receiving the data.

