

Introduction to MPI

March 20, 2015

What is MPI

- ▶ MPI is a library for passing messages between processes.
- ▶ MPI is a standard, not an implementation. We happen to be using the `openmpi` implementation.
- ▶ Since MPI is a standard, any correctly implemented code relying on MPI should be able to run using any properly implemented MPI library.
- ▶ For us, TORQUE is the mechanism for reserving resources. MPI is the mechanism that is helping us run processes on the those resources and communicate between the processes.

Basic concepts in MPI

- ▶ The fundamental tasks that MPI addresses are **sending** and **receiving** messages.
- ▶ MPI allows us to group processes
- ▶ MPI allows us to tag messages

- ▶ In this course we will use the **Single Program Multiple Data** model (SPMD). In this model all the processes are copies of the same executable.
- ▶ Behavior of each process is determined by the **rank**, which is managed by MPI.

MPI functions

The following functions make up most of the functionality of MPI:

- ▶ `MPI_Init` – initialize MPI
- ▶ `MPI_Comm_size` – get the number of processes in the given communicator
- ▶ `MPI_Comm_rank` – get the rank of this process in the given communicator
- ▶ `MPI_Send` – send a message
- ▶ `MPI_Recv` – receive a message
- ▶ `MPI_Finalize` – clean up

Hello, World

```
#include<mpi.h>
#include<stdio.h>

int main(int argc, char** argv){
    MPI_Init(&argc,&argv);

    printf("Hello, World\n");
    MPI_Finalize();

}
```

- ▶ To compile the previous code, use
`mpicc helloworld.c -o helloworld`
- ▶ To run the code, use
`mpirun helloworld` INSIDE AN INTERACTIVE QSUB
SESSION!

- ▶ Note also that `printf` worked properly, even though it was being executed on a remote node. This is the MPI runtime at work for us!
- ▶ `MPI_COMM_WORLD` is the world communicator, every process is included in this communicator.

We will frequently make behavior depend upon rank. The process with rank 0 is usually designated the “master” or “head”, and all the other processes are called “workers”.

```
#include<mpi.h>
#include<stdio.h>

int main(int argc, char** argv){
    MPI_Init(&argc,&argv);
    int total_procs;
    int rank;
    MPI_Comm_size(MPI_COMM_WORLD,&total_procs);
    MPI_Comm_rank(MPI_COMM_WORLD,&rank);
    if (rank==0){
        printf("Hello, this is process %d  and
I'm the master.\n",rank);
    }
    else{
        printf("Hello, this is process %d  and
I'm a lowly worker. :(\n",rank);
    }
    MPI_Finalize();
}
```

Sending and Receiving

- ▶ `int MPI_Send(void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm)`
 - ▶ `buf` is a pointer to the beginning of the data to send. Note that it can be of any type.
 - ▶ `count` is the number of data pieces to send
 - ▶ `datatype` is the type of data being sent (tells us how big the data pieces are), and is `MPI_INT`, `MPI_DOUBLE`, `MPI_CHAR`, etc...
 - ▶ `dest` is the rank of the process we are sending the message to
 - ▶ `tag` we will usually set to 0
 - ▶ `comm` we will set to `MPI_COMM_WORLD`

- ▶ `int MPI_Recv(void *buf, int count, MPI_Datatype datatype, int source, int tag, MPI_Comm comm, MPI_Status *status)`
 - ▶ `buf` is a pointer to the memory allocated to store the message
 - ▶ `count` is the size of the buffer, we can not receive more data than this.
 - ▶ `datatype` is the type of data we are receiving and is one of `MPI_INT`, `MPI_DOUBLE`, `MPI_CHAR`, etc
 - ▶ `source` is the rank of the process we want to receive from. Set to `MPI_ANY_SOURCE` to receive messages from anyone.
 - ▶ `tag` is the message tag we want to receive. Usually set to `MPI_ANY_TAG`
 - ▶ `comm` is always `MPI_COMM_WORLD`
 - ▶ `status` gets further information about the message
 - ▶ `status.MPI_SOURCE` tells us who sent the message
 - ▶ `MPI_Get_count(status)` tells us how long the message was

- ▶ Page 67 (Section 3.5 of the book) has a **quick** overview of MPI.
- ▶ <http://www.mpitutorial.com/>
- ▶ <http://www.lam-mpi.org/tutorials/one-step/ezstart.php>
- ▶ <http://www.mcs.anl.gov/research/projects/mpi/tutorial/index.html>

```
int main(int argc, char** argv){
    MPI_Init(&argc,&argv);
    int total_procs;
    int rank;
    MPI_Comm_size(MPI_COMM_WORLD,&total_procs);
    MPI_Comm_rank(MPI_COMM_WORLD,&rank);
    if (rank==0){
        int i;
        char* message="I am your master!";
        char recv_buffer[10];
        MPI_Status status;
        for (i=1;i<total_procs;i++){
            MPI_Send(message,18,MPI_CHAR,i,0,MPI_COMM_WORLD);
        }
        for (i=1;i<total_procs;i++){
            MPI_Recv(recv_buffer,10,MPI_CHAR,MPI_ANY_SOURCE,MPI_A
            printf("Process %d: Message from %d: %s\n",rank,status
        }
    }
}
```